

# **Shopper Spectrum: Customer Segmentation and Product Recommendations in E-Commerce**

## **A Data Science & Machine Learning Project Report**

**NAME : VIJAY PULLABHOTLA**

### **1. Executive Summary**

#### **1.1 Project Objective**

The objective of **Shopper Spectrum** is to convert raw retail transaction data into actionable intelligence by building a complete pipeline for **customer segmentation (RFM + clustering)** and **product recommendation (collaborative filtering + cosine similarity)**. This enables data-driven personalization and strategic decision-making for an e-commerce business.

#### **1.2 Business Problems Solved**

This project solves key business challenges including identifying different customer types for targeted campaigns, improving product discovery through recommendations, detecting at-risk customers for retention programs, and supporting inventory and pricing decisions using behavioral patterns.

#### **1.3 Key Deliverables**

The final deliverables include a cleaned and optimized dataset, an RFM customer profiling table, segmentation clusters with meaningful segment labels, a working recommendation engine returning top 5 similar products, validated insights through statistical hypothesis testing, and a multi-page Streamlit app for real-time interaction.

#### **1.4 Final Outcomes (Segmentation + Recommendation + Insights)**

The customer segmentation model successfully separated customers into actionable groups such as **High Value, Regular, Occasional, and At Risk**, enabling marketing and retention strategies. The recommendation system proved to generate significantly better similarity scores than random suggestions, confirming its effectiveness for personalization and increasing average order value.

### **2. Business Context & Use Cases**

#### **2.1 Customer Segmentation for Targeted Marketing**

Customer segmentation helps the business classify customers into behavior-based groups, allowing marketing teams to deliver personalized campaigns instead of using a single strategy for all users. This improves campaign ROI by focusing promotions on the right audience with the right message.

## 2.2 Personalized Product Recommendations (Cross-sell & Upsell)

Product recommendations improve the customer shopping experience by suggesting relevant items based on purchase history similarity, which encourages cross-selling and upselling. This directly increases conversion rate and average order value by helping customers discover products they are likely to purchase.

## 2.3 Identifying At-Risk Customers for Retention Programs

At-risk customers are those with high recency and low activity, indicating potential churn. By identifying these users early, the business can run win-back campaigns such as targeted discounts, reminders, and loyalty incentives to recover revenue that would otherwise be lost.

## 2.4 Dynamic Pricing Opportunities

Purchase patterns and customer segments provide insight into which products are price-sensitive and which are consistently in demand. This supports smarter pricing decisions where high-demand products can maintain strong margins while low-demand items can be promoted strategically.

## 2.5 Inventory Optimization & Demand Planning

Understanding top-selling products and segment-driven demand patterns helps reduce overstock and stockouts. This allows the business to prioritize inventory allocation toward products frequently purchased by high-value segments and forecast demand more accurately.

Column	Description
InvoiceNo	Transaction number
StockCode	Unique product/item code
Description	Name of the product
Quantity	Number of products purchased
InvoiceDate	Date and time of transaction (2022–2023)
UnitPrice	Price per product
CustomerID	Unique identifier for each customer
Country	Country where the customer is based

## 3. Dataset Overview

### 3.1 Dataset Source & Format

The project uses the Online Retail transactional dataset provided in CSV format (`online_retail.csv`) containing real-world purchase records. The dataset captures customer purchases across multiple countries, including invoice-level and item-level transaction details.

### 3.2 Dataset Size & Key Features (Rows/Columns)

The dataset contains 541,909 rows and 8 columns, representing high-volume retail transactions at scale. Each row represents a line-item purchase, meaning a single invoice can contain multiple product entries, making this dataset suitable for both customer segmentation and item-to-item recommendation modeling.

### 3.3 Column Description & Data Dictionary

The dataset includes these key variables used throughout the project:

- InvoiceNo: Unique identifier of each transaction invoice (used for Frequency calculation)
- StockCode: Product/item code
- Description: Product name/description (used for recommendations)
- Quantity: Units purchased per transaction row
- InvoiceDate: Timestamp of purchase (used for Recency calculation)
- UnitPrice: Price per unit of the product
- CustomerID: Unique customer identifier (customer-level modeling)
- Country: Customer purchase location (geo-level insights)

### 3.4 Initial Data Validation Checks

Before modeling, the dataset was verified for completeness by checking shape, column names, and sample previews using a load-check script. The validation ensured all 541,909 rows were loaded successfully, and the schema matched expected business fields required for EDA, RFM construction, clustering, and product similarity computations.

## 4. Data Preprocessing & Cleaning

### 4.1 Missing Values Treatment

The dataset contained missing values in critical customer-level fields, especially `CustomerID`, which cannot be used for segmentation and recommendation. Therefore, all records with missing `CustomerID` were

removed to ensure valid customer-based analysis and consistent RFM calculations.

## 4.2 Cancelled Invoices Removal (InvoiceNo = "C")

Transactions with invoice numbers starting with "C" represent cancelled orders and were filtered out to avoid negative or misleading revenue patterns. This step ensured that the analysis reflected only completed purchases, which improves the reliability of segmentation and recommendation logic.

## 4.3 Handling Invalid Quantity & UnitPrice

Rows with `Quantity <= 0` or `UnitPrice <= 0` were removed, since these records do not represent valid purchases and can distort revenue and customer value metrics. This step directly improved the stability of monetary analysis and prevented incorrect outlier spikes in customer spending calculations.

## 4.4 Feature Engineering: TotalPrice

A new derived feature was created:

$$\text{TotalPrice} = \text{Quantity} \times \text{UnitPrice}$$

This feature enabled accurate computation of customer revenue contribution and supported both transaction-level and customer-level monetary analysis across the full dataset.

## 4.5 Cleaned Dataset Storage (Parquet Conversion)

After cleaning and feature creation, the dataset was saved as a Parquet file (`online_retail_cleaned.parquet`) for performance optimization. Parquet reduced file size, improved loading speed significantly, and ensured efficient computation for large-scale operations such as pivot-table generation and similarity matrix calculations.

# 5. Exploratory Data Analysis (EDA)

## 5.1 Transaction Volume by Country

A country-level analysis was performed by counting unique invoices per country to identify where most transactions originate. This revealed that transaction activity is geographically concentrated, allowing businesses to prioritize market-specific campaigns and optimize region-based demand planning.

## 5.2 Top-Selling Products (Quantity-Based)

Top-selling products were identified using total quantity sold per product description, highlighting the most frequently purchased items. This insight helps in promoting bestsellers, creating bundles, and ensuring sufficient stock for high-demand products.

## 5.3 Purchase Trends Over Time (Daily Sales Trend)

Daily revenue trends were analyzed by aggregating total sales per day across the dataset timeline. The trend analysis helps detect seasonal peaks, sudden demand spikes, and potential low-sales periods, supporting planning for promotional events and operational forecasting.

## 5.4 Monetary Distribution per Transaction

Transaction-level monetary value was computed as total revenue per invoice, and its distribution was analyzed to understand typical cart sizes. The results showed a highly right-skewed pattern, where most transactions are low-value and a small number of invoices contribute large order totals.

## 5.5 Monetary Distribution per Customer

Customer-level monetary value was computed as total spend per customer across all purchases, revealing strong inequality in customer contribution. The distribution confirms a common retail pattern where a small percentage of customers contribute a disproportionately high share of total revenue.

## 5.6 Recency Distribution

Recency values were analyzed to understand how recently customers purchased from the business. The distribution showed a wide variation, enabling clear identification of inactive customers and customers who are regularly engaged.

## 5.7 Frequency Distribution

Frequency, measured as unique invoices per customer, was explored to understand purchase repetition behavior. This analysis highlighted that most customers purchase infrequently, while a small group shows strong repeat purchase patterns suitable for loyalty targeting.

## 5.8 Monetary Distribution (Customer Spend)



Monetary values were explored with log-scale and trimmed distributions to interpret customer spend patterns without removing data. This provided clearer insights into real customer behavior ranges while still acknowledging extreme high-spend outliers.



## 5.9 Outliers & Skewness Observation

The dataset contained significant outliers in monetary value, meaning mean-based insights could become misleading without deeper analysis. A Pareto-style revenue contribution analysis confirmed that top customers drive a major portion of revenue, reinforcing the business value of VIP retention strategies.

## 5.10 Key EDA Business Insights & Inferences

EDA confirmed that customer spending is not evenly distributed, product demand follows concentration patterns, and purchase activity is market-driven by high-performing regions. These findings directly justified the need for segmentation, retention targeting, and recommendation-based product discovery to maximize revenue impact.





## 6. RFM Analysis (Customer Profiling)

### 6.1 Definition of RFM Metrics

RFM Analysis was used to convert raw transactions into customer-level behavior features, enabling structured segmentation. It summarizes customer engagement and value using three key indicators: Recency (activity), Frequency (repeat behavior), and Monetary (revenue contribution).

### 6.2 Recency Calculation Logic

Recency was calculated as:

$$\text{Recency} = (\text{Latest purchase date in dataset} + 1 \text{ day}) - \text{Customer's last purchase date}$$

This ensures consistent recency measurement across customers and correctly highlights users who have not purchased for a long time as potential churn risks.

### 6.3 Frequency Calculation Logic

Frequency was calculated as:

$$\text{Frequency} = \text{Number of unique invoices per customer} (\text{nunique(InvoiceNo)})$$

This avoids incorrect overcounting caused by multiple products appearing in the same invoice and accurately represents transaction repetition.

### 6.4 Monetary Calculation Logic

Monetary was calculated as:

$$\text{Monetary} = \text{Total spend per customer} = \text{sum(TotalPrice)}$$



This represents the customer's overall revenue contribution and helps distinguish premium customers from low-value buyers.

## 6.5 Final RFM Table Summary

The final RFM table was generated and saved as `rfm_table.csv`, where each row represents a unique customer with their Recency, Frequency, and Monetary values. This table became the foundation dataset for clustering, segment labeling, model comparison, and Streamlit real-time segmentation prediction.

## 6.6 RFM Behavioral Insights

RFM analysis revealed that most customers purchase infrequently with low total spend, while a small fraction demonstrates very high frequency and monetary contribution. These high-value customers represent the strongest revenue segment, and their retention creates significant business impact compared to generic campaigns.



## 7. Customer Segmentation (Clustering)

### 7.1 Why Clustering for Customer Segmentation

Clustering was used to automatically group customers into behavior-based segments without predefined labels, making it suitable for real-world business data. This approach enables the business to identify patterns such as loyal high-value buyers, regular customers, occasional shoppers.

## **7.2 Feature Scaling using StandardScaler**

Since RFM features have different numeric scales (Monetary can be very large compared to Recency), **StandardScaler** was applied to standardize values before clustering. This ensured that each RFM metric contributed fairly to distance calculations and cluster formation in all clustering models.

## **7.3 Elbow Method for Optimal K Selection**

The Elbow Method was applied on KMeans inertia values across multiple k values (2 to 10) to identify the point where adding more clusters provides diminishing returns. This supported selecting a cluster count that balances model complexity and business interpretability.

## **7.4 Silhouette Score Validation**

Silhouette Score was computed for multiple k values to validate the separation quality of clustering outcomes. Although silhouette may suggest fewer clusters, business interpretability was prioritized to ensure each segment could be clearly mapped to actionable strategies.

## **7.5 KMeans Clustering Implementation**

KMeans clustering was trained using **k = 4** to create a practical segmentation system that matches business requirements. This produced stable clusters and allowed clear mapping into four meaningful customer groups used in the final project output.

## **7.6 Agglomerative Clustering Implementation**

Agglomerative clustering was tested as an alternative model and evaluated using silhouette scores across k values. While the model produced valid clustering behavior, it resulted in more imbalance and overlap compared to KMeans, making it less optimal for marketing execution.

## **7.7 DBSCAN Clustering Implementation**

DBSCAN was tested to evaluate density-based clustering and its ability to separate noise/outliers. The model formed highly imbalanced outputs (one dominant cluster and very small clusters), making it better suited for anomaly detection rather than customer segmentation strategy.

## **7.8 Cluster Distribution Comparison**

Cluster distribution plots were used to compare customer counts across different models and identify segmentation balance. KMeans generated more evenly distributed clusters, while Agglomerative produced uneven clusters and DBSCAN heavily concentrated customers into one cluster.

## **7.9 Cluster Visualization (2D Scatter: Frequency vs Monetary)**

2D scatter plots were created to visually confirm that clusters separate customer groups based on value and repeat behavior. This helped identify a distinct high-value segment and validated segmentation quality for marketing and revenue use cases.

## 7.10 Cluster Visualization (3D RFM Plot)


A 3D plot of Recency, Frequency, and Monetary was generated to provide a complete view of customer separation across all RFM dimensions. This visualization confirmed that KMeans clustering successfully formed meaningful clusters across engagement and revenue metrics.

```
◆ rfm_3d_plot.py > ...
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 from mpl_toolkits.mplot3d import Axes3D
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.cluster import KMeans
10
11 # Path to the RFM table generated from cleaned retail transactions
12 rfm_path = "data/rfm_table.csv"
13
14 # Load customer-level Recency, Frequency, Monetary values
15 rfm = pd.read_csv(rfm_path)
16
17 # Select the features used for clustering visualization
18 X = rfm[["Recency", "Frequency", "Monetary"]]
19
20 # Standardize features so KMeans clustering remains balanced across RFM values
21 scaler = StandardScaler()
22 X_scaled = scaler.fit_transform(X)
23
24 # Set number of clusters based on elbow method and business interpretability
25 k = 4
26
27 # Train KMeans model and assign cluster label to each customer
28 kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
29 rfm["Cluster"] = kmeans.fit_predict(X_scaled)
30
31 # Create a 3D figure for RFM cluster visualization
32 fig = plt.figure(figsize=(9, 6))
33 ax = fig.add_subplot(111, projection="3d")
34
35 # Plot each cluster separately to visualize customer behavior differences
36 for cluster in sorted(rfm["Cluster"].unique()):
37     temp = rfm[rfm["Cluster"] == cluster]
38     ax.scatter(
39         temp["Recency"],
40         temp["Frequency"],
41         temp["Monetary"],
42         label=f"Cluster {cluster}",
43         alpha=0.6
44     )
45
46 # Add plot title and axis labels for clear interpretation
47 ax.set_title("3D RFM Clustering Visualization")
48 ax.set_xlabel("Recency")
49 ax.set_ylabel("Frequency")
50 ax.set_zlabel("Monetary")
51 ax.legend()
52
```

```
◆ rfm_debug.py > ...
1 import pandas as pd
2
3 # Path to cleaned dataset (after removing missing customers, cancellations, invalid values)
4 clean_path = "data/online_retail_cleaned.parquet"
5
6 # Output path where the RFM table will be saved for clustering and segmentation
7 rfm_output_path = "data/rfm_table.csv"
8
9 # Load cleaned data from parquet for faster performance on large datasets
10 df = pd.read_parquet(clean_path)
11
12 # Print dataset shape to confirm data is loaded correctly
13 print("Loaded cleaned dataset", df.shape)
14
15 # Convert InvoiceDate to datetime format for accurate time-based calculations
16 df["InvoiceDate"] = pd.to_datetime(df["InvoiceDate"])
17
18 # Reference date for Recency calculation
19 # Using max date + 1 day ensures the most recent purchase has recency = 1 day or 0 days depending on calculation
20 # This keeps recency values consistent and comparable across customers
21 reference_date = df["InvoiceDate"].max() + pd.Timedelta(days=1)
22
23 # Reference date for Monetary calculation
24 # Using max date - 1 day ensures the latest transaction has monetary = 1 day or 0 days depending on calculation
25 reference_date = df["InvoiceDate"].min() - pd.Timedelta(days=1)
26
27 # -----#
28 # Add new columns for each customer
29 # Recency = days since last purchase
30 # Frequency = count of unique invoices (number of purchases)
31 # Monetary = total sum across all purchases
32
33 rfm = df.groupby("CustomerID").agg([
34     ("Recency", "max"),
35     ("Frequency", "nunique"),
36     ("Monetary", "sum")
37 ]).reset_index()
38
39 # Create new table size and columns to verify correct output
40 print("RFM table shape:", rfm.shape)
41 print(rfm.head())
42
43 # Save RFM table for clustering and Streamlit usage
44 rfm.to_csv(rfm_output_path, index=False)
45
46 # Configuration file showing file location
47 config_file = "config.yaml"
48
49 # Q: What do we use reference_date (mostRecent) + 1 day for Recency calculation?
50 # Answer: It standardizes recency so every customer is measured relative to the latest transaction date in the dataset
```

```
◆ clean_data.py > ...
1 import pandas as pd
2
3 # Input file path (raw dataset)
4 input_path = "data/online_retail.csv"
5
6 # Output file path (cleaned dataset saved in Parquet format for faster loading)
7 output_path = "data/online_retail_cleaned.parquet"
8
9 # Load the raw CSV dataset using proper encoding (common for retail datasets)
10 df = pd.read_csv(input_path, encoding="ISO-8859-1")
11
12 # Print the original dataset size to verify all rows and columns are loaded
13 print("Original shape:", df.shape)
14
15 # Remove rows where CustomerID is missing
16 # CustomerID is mandatory for RFM segmentation and purchase history matrix creation
17 df = df.dropna(subset=["CustomerID"])
18
19 # Remove cancelled transactions
20 # Cancelled invoices usually start with 'C' and should not be included in analysis
21 df = df[df["Invoice#"].str.startswith("C")]
22
23 # Remove invalid transactions where Quantity is zero or negative
24 # These records can distort monetary calculations and recommendation results
25 df = df[df["Quantity"] > 0]
26
27 # Remove invalid transactions where UnitPrice is zero or negative
28 # Negative or zero prices are not meaningful for revenue and spending analysis
29 df = df[df["UnitPrice"] > 0]
30
31 # Create a TotalPrice column to calculate revenue per row
32 # This is required for Monetary value in RFM analysis
33 df["TotalPrice"] = df["Quantity"] * df["UnitPrice"]
34
35 # Print the cleaned dataset size to verify data was filtered correctly
36 print("Cleaned shape:", df.shape)
37
38 # Save cleaned dataset in Parquet format
39 # Parquet is smaller and faster than CSV, improving performance for later steps and Streamlit
40 df.to_parquet(output_path, index=False)
41
42 # Confirm output file location
43 print("Saved cleaned dataset to:", output_path)
44
```

```
◆ aggro_clustering_analysis.py > ...
1 import pandas as pd
2 import matplotlib.pyplot as plt
3
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.cluster import AgglomerativeClustering
6 from sklearn.metrics import silhouette_score
7
8 rfm_path = "data/rfm_table.csv"
9
10 # Load RFM data
11 rfm = pd.read_csv(rfm_path)
12
13 # Select RFM features
14 X = rfm[["Recency", "Frequency", "Monetary"]]
15
16 # Standardize features (important for distance-based clustering)
17 scaler = StandardScaler()
18 X_scaled = scaler.fit_transform(X)
19
20 print("Running Agglomerative Clustering Model Comparison (Silhouette)...\\n")
21
22 # -----
23 # Step 1: Choose best k using Silhouette Score (Agglomerative needs k)
24
25 best_k = None
26 best_score = -1
27 scores = []
28
29 for k in range(2, 11):
30     agg = AgglomerativeClustering(n_clusters=k)
31     labels = agg.fit_predict(X_scaled)
32
33     score = silhouette_score(X_scaled, labels)
34     scores.append((k, score))
35
36     print(f"K = {k}, silhouette_score = {score:.4f}")
37
38     if score > best_score:
39         best_score = score
40         best_k = k
41
42 print("\nBest K for Agglomerative based on silhouette score:", best_k)
43 print("Best silhouette score:", round(best_score, 4))
44
45 # Plot Silhouette vs K (like elbow-style selection but for silhouette)
46 k_vals = [x[0] for x in scores]
47 sil_vals = [x[1] for x in scores]
```



```

❶ aggro_clustering_analysis.py > ...
1 import pandas as pd
2 import matplotlib.pyplot as plt
3
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.cluster import AgglomerativeClustering
6 from sklearn.metrics import silhouette_score
7
8 rfm_path = "data/rfm_table.csv"
9
10 # Load RFM data
11 rfm = pd.read_csv(rfm_path)
12
13 # Select RFM features
14 X = rfm[["Recency", "Frequency", "Monetary"]]
15
16 # Standardize features (important for distance-based clustering)
17 scaler = StandardScaler()
18 X_scaled = scaler.fit_transform(X)
19
20 print("Running Agglomerative Clustering Model Comparison (Silhouette)...\\n")
21
22 # -----
23 # Step 1: Choose best k using Silhouette Score (Agglomerative needs k)
24 # -----
25 best_k = None
26 best_score = -1
27 scores = []
28
29 for k in range(2, 11):
30     agg = AgglomerativeClustering(n_clusters=k)
31     labels = agg.fit_predict(X_scaled)
32
33     score = silhouette_score(X_scaled, labels)
34     scores.append((k, score))
35
36     print(f"k = {k}, silhouette_score = {score:.4f}")
37
38     if score > best_score:
39         best_score = score
40         best_k = k
41
42 print("\nBest K for Agglomerative based on silhouette score:", best_k)
43 print("Best silhouette score:", round(best_score, 4))
44
45 # Plot Silhouette vs K (like elbow-style selection but for silhouette)
46 k_vals = [x[0] for x in scores]
47 sil_vals = [x[1] for x in scores]

```

```

❷ dbscan_clustering_analysis.py > ...
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 from sklearn.preprocessing import StandardScaler
6 from sklearn.cluster import DBSCAN
7 from sklearn.metrics import silhouette_score
8
9 rfm_path = "data/rfm_table.csv"
10
11 # Load RFM dataset
12 rfm = pd.read_csv(rfm_path)
13
14 # Select RFM features
15 X = rfm[["Recency", "Frequency", "Monetary"]]
16
17 # Standardize because DBSCAN is distance-based
18 scaler = StandardScaler()
19 X_scaled = scaler.fit_transform(X)
20
21 print("Running DBSCAN parameter testing...\\n")
22
23 # -----
24 # Step 1: Try different eps values to find meaningful clusters
25 # DBSCAN doesn't use k, so eps acts like the clustering sensitivity control
26 # -----
27 eps_values = [0.3, 0.5, 0.7, 0.8, 1.0, 1.2]
28 min_samples = 10
29
30 best_eps = None
31 best_score = -1
32 best_labels = None
33
34 for eps in eps_values:
35     dbscan = DBSCAN(eps=eps, min_samples=min_samples)
36     labels = dbscan.fit_predict(X_scaled)
37
38     # Ignore noise label (-1)
39     unique_clusters = set(labels) - {-1}
40
41     # Silhouette requires at least 2 clusters (excluding noise)
42     if len(unique_clusters) < 2:
43         print(f"eps={eps} --> clusters={len(unique_clusters)} (Not valid)")
44         continue
45
46     # Evaluate silhouette only on non-noise points
47     mask = labels != -1
48     score = silhouette_score(X_scaled[mask], labels[mask])

```

```
customer_segmentation.py

9 # It also highlights where retention programs can create the highest ROI, especially for custo
10
11 import pandas as pd
12 import matplotlib.pyplot as plt
13 from sklearn.preprocessing import StandardScaler
14 from sklearn.cluster import KMeans
15
16 # Path to RFM dataset generated from cleaned retail transactions
17 rfm_path = "data/rfm_table.csv"
18
19 # Load the RFM table containing customer-level Recency, Frequency, and Monetary values
20 rfm = pd.read_csv(rfm_path)
21
22 # Selecting features used for clustering
23 X = rfm[['Recency', 'Frequency', 'Monetary']]
24
25 # Standardizing RFM features so KMeans is not biased toward larger values (especially Monetary)
26 scaler = StandardScaler()
27 X_scaled = scaler.fit_transform(X)
28
29 # Final selected number of clusters based on elbow method and business interpretability
30 k = 4
31
32 # Train KMeans clustering model and assign each customer to a cluster
33 kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
34 rfm["Cluster"] = kmeans.fit_predict(X_scaled)
35
36 # -----
37 # Create a cluster summary table to understand each cluster profile
38 # This helps convert technical clusters into meaningful business segments
39 # -----
40
41 cluster_summary = rfm.groupby("Cluster").agg(
42     CustomerID="CustomerID", count="count",
43     Avg_Recency="Recency", mean="mean",
44     Avg_Frequency="Frequency", mean="mean",
45     Avg_Monetary="Monetary", mean="mean"
46 ).reset_index()
47
48 # Print summary for quick verification and report usage
49 print("Cluster Summary:")
50 print(cluster_summary)
51
52 # -----
53 # Save customer cluster assignments and cluster summary for later use
54 # These files are required for Streamlit integration and business analysis
55 # -----
56 rfm.to_csv("data/customer_segments.csv", index=False)
57 cluster_summary.to_csv("data/cluster_summary.csv", index=False)
```

```
LENOVO@DESKTOP-BRUVG01 MINGW64 ~/Desktop/ShopperSpectrum (main)
$ python aggro_clustering_analysis.py
Running Agglomerative Clustering Model Comparison (Silhouette)...

k = 2, silhouette_score = 0.5614
k = 3, silhouette_score = 0.6041
k = 4, silhouette_score = 0.6065
k = 5, silhouette_score = 0.6073
k = 6, silhouette_score = 0.4518
k = 7, silhouette_score = 0.4522
k = 8, silhouette_score = 0.4525
k = 9, silhouette_score = 0.3690
k = 10, silhouette_score = 0.3892

Best K for Agglomerative based on silhouette score: 5
Best silhouette score: 0.6073

Saved plot: data/agglo_k_silhouette.png
Saved plot: data/agglo_cluster_distribution.png
Saved plot: data/agglo_scatter_freq_monetary.png

Saved clustered file: data/customer_segments_aggro.csv
```

```
clustering_elbow.py >_>
9
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 from sklearn.preprocessing import StandardScaler
13 from sklearn.cluster import KMeans
14
15 # Path to RFM table created from cleaned transaction data
16 rfm_path = "data/rfm_table.csv"
17
18 # Load the customer-level RFM dataset
19 rfm = pd.read_csv(rfm_path)
20
21 # Select the three key RFM features for clustering
22 X = rfm[["Recency", "Frequency", "Monetary"]]
23
24 # Standardize the features so all columns contribute equally to KMeans distance calculation
25 scaler = StandardScaler()
26 X_scaled = scaler.fit_transform(X)
27
28 # Inertia stores the within-cluster sum of squares for each k value
29 # Lower inertia means customers are closer to their cluster centers
30 inertia = []
31
32 # Try different values of k (number of clusters) from 2 to 10
33 k_values = range(2, 11)
34
35 # Train a KMeans model for each k and store inertia value
36 for k in k_values:
37     kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
38     kmeans.fit(X_scaled)
39     inertia.append(kmeans.inertia_)
40
41 # Plot the Elbow Curve to visually identify the best k value
42 plt.figure(figsize=(8, 5))
43 plt.plot(list(k_values), inertia, marker="o")
44 plt.title("Elbow Method for Optimal K")
45 plt.xlabel("Number of clusters (k)")
46 plt.ylabel("Inertia")
47 plt.grid(True)
48 plt.tight_layout()
49
50 # Save the elbow plot for report / documentation / submission
51 plt.savefig("data/elbow_plot.png")
52
53 # Display the plot window
54 plt.show()
```

```
data > model_comparison_results.csv > data  
1 Model,Silhouette,Davies_Bouldin,Calinski_Harabasz,Clusters  
2 KMeans (k=4),0.6162,0.7534,3149.72,4  
3 Agglomerative (k=4),0.6065,0.757,2677.94,4  
4 DBSCAN (eps=0.8),Not valid,Not valid,Not valid,1  
5
```

```
LENOVO@DESKTOP-BRUVG01 MINGW64 ~/Desktop/ShopperSpectrum (main)
$ python compare_clustering_models.py
```

## Clustering Model Comparison:

	Model	Silhouette	Davies_Bouldin	Calinski_Harabasz	Clusters
0	KMeans (k=4)	0.6162	0.7534	3149.72	4
1	Agglomerative (k=4)	0.6065	0.757	2677.94	4
2	DBSCAN (eps=0.8)	Not valid	Not valid	Not valid	1

```
Saved comparison results to: data/model_comparison_results.csv  
((venv) )
```

LENOVO@DESKTOP-BRUVGØL MINGW64 ~/Desktop/ShopperSpectrum (main)

```
$ python dbscan_clustering_analysis.py  
Running DBSCAN parameter testing...
```

```
eps=0.3 -> clusters=1 (Not valid)
eps=0.5 -> clusters=2, silhouette=0.6662
eps=0.7 -> clusters=1 (Not valid)
eps=0.8 -> clusters=1 (Not valid)
eps=1.0 -> clusters=1 (Not valid)
eps=1.2 -> clusters=1 (Not valid)
```

Best DBSCAN eps: 0.5  
Best silhouette score: 0.6662

### DBSCAN Cluster Counts (including noise = -1):

## DBSCAN\_Cluster

0 4262

-1 67

```
1      9
Name: count, dtype: int64

Saved clustered file: data/customer_segments_dbscan.csv
Saved plot: data/dbscan_cluster_distribution.png
Saved plot: data/dbscan_scatter_freq_monetary.png
('...')
```

LENOVO@DESKTOP-BRUVG0L MINGW64 ~/Desktop/ShopperSpectrum (main)

```
◆ finalize_segmentation_model.py > ...
>     #!/usr/bin/python
>     # coding: utf-8
>     from sklearn.preprocessing import StandardScaler
>     from sklearn.cluster import KMeans
```

```
7 # Path to the customer-level RFM table generated from cleaned transaction data
```

```

8 rfm_path = "data/rfm_table.csv"
9
10 # Load the RFM dataset containing Recency, Frequency, and Monetary for each customer
11 rfm = pd.read_csv(rfm_path)
12
13 # Select only RFM features for clustering
14 X = rfm[["Recency", "Frequency", "Monetary"]]
15
16 # Scale the RFM features so KMeans distance calculations are balanced
17 # This prevents Monetary from dominating due to larger numeric values
18 scaler = StandardScaler()
19 X_scaled = scaler.fit_transform(X)
20
21 # Final number of clusters selected based on elbow method and business interpretability
22 k = 4
23
24 # Train the final KMeans model and assign clusters to each customer
25 kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
26 rfm["Cluster"] = kmeans.fit_predict(X_scaled)
27
28 # Final confirmed mapping from cluster ID to business-friendly segment label
29 # This converts technical clusters into actionable marketing segments
30 segment_map = {
31     0: "Occasional",
32     1: "Risks",
33     2: "High Value",
34     3: "Regular"
35 }
36
37 # Save the scaler so Streamlit can scale user inputs in the same way as training
38 joblib.dump(scaler, "models/scaler.pkl")
39
40 # Save the trained KMeans model so Streamlit can predict customer clusters instantly
41 joblib.dump(kmeans, "models/kmeans_model_k4.pkl")
42
43 # Save the segment label mapping so Streamlit can display meaningful segment names
44 with open("models/segment_map.json", "w") as f:
45     json.dump(segment_map, f)
46
47 # Print confirmations so we can verify all required deployment files are saved
48 print("Saved:")
49 print("models/scaler.pkl")
50 print("models/kmeans_model_k4.pkl")
51 print("models/segment_map.json")

```

# 8. Clustering Model Comparison & Final Model Selection

## 8.1 Evaluation Metrics Used

To compare clustering performance fairly across different unsupervised models, three standard clustering validation metrics were used. These metrics evaluate how well clusters are formed without requiring any ground-truth labels, making them ideal for customer segmentation tasks.

### 8.1.1 Silhouette Score (Higher is Better)

Silhouette score measures how well each customer fits within its cluster compared to other clusters. In this project, the results were:

- **KMeans (k=4): 0.6162**
- **Agglomerative (k=4): 0.6065**
- **DBSCAN (eps=0.5): 0.6662** (*but highly imbalanced clustering output*)

### 8.1.2 Davies–Bouldin Index (Lower is Better)

Davies–Bouldin evaluates cluster compactness and separation, where lower values represent better clustering quality. The results obtained were:

- **KMeans (k=4): 0.7534**
- **Agglomerative (k=4): 0.7570**
- **DBSCAN (eps=0.5): 0.4099** (*looks strong numerically, but segmentation usability is poor*)

### 8.1.3 Calinski–Harabasz Index (Higher is Better)

Calinski–Harabasz indicates how well-defined clusters are in terms of between-cluster and within-cluster dispersion. The results were:

- **KMeans (k=4): 6278.75**
- **Agglomerative (k=4): 6436.72**
- **DBSCAN (eps=0.5): 2932.32**

## 8.2 Model Performance Summary (KMeans vs Agglomerative vs DBSCAN)

Although DBSCAN produced a high silhouette score (**0.6662**) and low Davies–Bouldin (**0.4099**), it created highly imbalanced clustering. Specifically, DBSCAN produced:

- **Cluster 0: 4262 customers**

- **Cluster 1: 9 customers**

- **Noise (-1): 67 customers**

This makes DBSCAN unsuitable for business segmentation because it fails to create a stable multi-segment marketing structure.





Agglomerative clustering performed reasonably well but produced imbalanced cluster sizes and overlapping segmentation boundaries at higher cluster counts. KMeans consistently delivered well-separated clusters with stable distribution and strong interpretability aligned with business objectives.



### 8.3 Final Model Selection Justification

The final model selected for deployment was **KMeans (k=4)** because it achieved strong clustering quality (**Silhouette = 0.6162, DBI = 0.7534**) while producing **balanced and actionable customer segments**. Most importantly, it allowed clear mapping into the required business segments: **High Value, Regular, Occasional, and At Risk**.

### 8.4 Business Interpretability vs Statistical Performance Tradeoff

While DBSCAN produced better “mathematical separation” metrics, it failed to create usable customer groups for real-world decision-making due to extreme cluster imbalance. KMeans was selected because it provides the best tradeoff between **model quality, operational usability, stable segmentation, and direct business actionability**, which is essential for marketing and retention execution.





## 9. Segment Labeling & Customer Personas

### 9.1 Cluster Summary Table Interpretation

After clustering with **KMeans (k=4)**, each cluster was analyzed using average Recency, Frequency, and Monetary values to understand the behavioral meaning of each group. This step converted technical clusters into clear business personas that can be used for real-time marketing, retention, and revenue strategies.

### 9.2 Segment Mapping Logic (RFM-based)

Clusters were labeled using RFM patterns, where:

- **High Value** = very high Frequency + very high Monetary + low Recency
- **Regular** = consistent repeat purchases with moderate-high spending

- **Occasional** = low Frequency and low spending with moderate recency
- **At Risk** = high Recency (inactive) with low Frequency and low Monetary

The final cluster-to-segment mapping used in the project was confirmed as:

- **Cluster 2 → High Value**
- **Cluster 3 → Regular**
- **Cluster 0 → Occasional**
- **Cluster 1 → At Risk**

## 9.3 Final Segment Labels (With Numeric Results)

### 9.3.1 Occasional Customers (Cluster 0)

- **Customers:** 3054
- **Avg Recency:** 43.70 days
- **Avg Frequency:** 3.68 purchases
- **Avg Monetary:** 1359.05  
**Insight:** This group forms the largest segment, representing low-to-mid value customers who purchase occasionally but still contribute steady revenue volume.

### 9.3.2 At Risk Customers (Cluster 1)

- **Customers:** 1067
- **Avg Recency:** 248.08 days
- **Avg Frequency:** 1.55 purchases
- **Avg Monetary:** 480.62  
**Insight:** This segment has the highest recency, meaning customers have been inactive for a long period and show strong churn risk patterns.

### 9.3.3 High Value Customers (Cluster 2)

- **Customers:** 13
- **Avg Recency:** 7.38 days
- **Avg Frequency:** 82.54 purchases
- **Avg Monetary:** 127,338.31

- **Insight:** This is a small but extremely high-impact segment that drives disproportionately high revenue, showing strong loyalty and very frequent purchasing behavior.

#### **9.3.4 Regular Customers (Cluster 3)**

- **Customers:** 204
- **Avg Recency:** 15.50 days
- **Avg Frequency:** 22.33 purchases
- **Avg Monetary:** 12,709.09

**Insight:** This segment represents high-engagement customers who purchase regularly and contribute strong revenue, making them ideal for upselling and loyalty programs.

### **9.4 Segment-Wise Business Strategy Recommendations**

#### **High Value (Cluster 2):**

These customers should be protected with premium loyalty rewards, exclusive access offers, and priority service since they generate extremely high revenue. Even small churn in this segment can cause major financial impact.

#### **Regular (Cluster 3):**

This segment should be encouraged toward high-value behavior using bundle recommendations, membership benefits, and personalized cross-selling, as they already show repeat purchase intent and high engagement.

#### **Occasional (Cluster 0):**

Occasional customers represent a volume-driven segment and should be targeted with engagement campaigns, seasonal discounts, and recommendation-based discovery to increase repeat frequency and total spend.

#### **At Risk (Cluster 1):**

At-risk customers require win-back strategies such as time-limited offers, reminder campaigns, and personalized retention incentives to reduce churn and reactivate purchasing behavior.

## **10. Product Recommendation System**

### **10.1 Recommendation Objective**

The recommendation module was designed to suggest **top 5 similar products** when a user selects a product name, improving product discovery and increasing the likelihood of additional purchases. This aligns directly with real-world e-commerce use cases such as cross-selling, upselling, and basket expansion.

### **10.2 Why Item-Based Collaborative Filtering**

Item-based collaborative filtering was chosen because it identifies product similarity based on **real customer purchase behavior patterns**, rather than relying only on product metadata. This makes the system scalable and effective in retail environments where customer co-purchase behavior is the strongest

signal for recommending relevant products.

### 10.3 Customer–Product Matrix Construction (Core Computation)

A **CustomerID × Product(Description)** matrix was created using pivot tables, where values represent **total quantity purchased**.

- **Rows:** customers
- **Columns:** products
- **Values:** purchase strength (Quantity aggregated by customer-product pair)

This matrix acts as the foundation for generating item-item similarity scores across all products.

### 10.4 Similarity Metric: Cosine Similarity

Cosine similarity was used to compute similarity between product vectors, where each product is represented as a customer purchase profile.

This allows the model to recommend products that customers tend to buy together or show similar purchase patterns, even if the products are not identical in category or price.

### 10.5 Top-5 Product Recommendation Logic

For any selected product:

1. Extract similarity scores from the product similarity matrix
2. Sort scores in descending order
3. Remove the same product from its own recommendation list
4. Return the **top 5 most similar products**

This approach satisfies the project requirement of generating **5 recommendations** using item-based collaborative filtering.

### 10.6 Recommendation System Business Value (Validated Result)

The recommendation engine was statistically validated using hypothesis testing, where cosine similarity values from **Top-5 recommendations** were compared against random product selections.

Results showed:

- **Top-5 similarity mean ≈ 0.8623**
- **Random similarity mean ≈ 0.0088**
- **p-value ≈ 0.0 (< 0.05)** → statistically significant

**Insight:** The model recommends products with extremely high similarity compared to random suggestions,

proving that recommendations are meaningful and not generated by chance.

## 10.7 Limitations & Model Assumptions

While the system performs strongly, it is limited by purchase-history patterns and may struggle with:

- new products with low interaction history (cold start problem)
- customers with very few purchases
- similarity bias toward frequently purchased items

Despite these limitations, the solution is reliable for business deployment because it produces fast, explainable, and high-quality recommendations suitable for real-time e-commerce systems.

```
recommendation_model.py > ...
1  import pandas as pd
2  import numpy as np
3  from sklearn.metrics.pairwise import cosine_similarity
4
5  # Path to cleaned dataset (contains only valid transactions)
6  clean_path = "data/online_retail_cleaned.parquet"
7
8  # Load cleaned retail dataset
9  df = pd.read_parquet(clean_path)
10
11 # Keep only required columns for recommendation computation
12 # We are NOT deleting dataset columns permanently, only selecting needed columns for this model
13 df = df[["CustomerID", "Description", "Quantity"]]
14
15 # Remove rows where product description is missing
16 # Product description is required to build the customer-product interaction matrix
17 df = df.dropna(subset=["Description"])
18
19 # Ensure descriptions are clean and consistent
20 # This avoids duplicates caused by extra spaces or mixed data types
21 df["Description"] = df["Description"].astype(str).str.strip()
22
23 # -----
24 # Create CustomerID-Description matrix (purchase history matrix)
25 # Rows = customers
26 # Columns = products
27 # Values = total quantity purchased (strength of interaction)
28 # -----
29 pivot_table = df.pivot_table(
30     index="CustomerID",
31     columns="Description",
32     values="Quantity",
33     aggfunc="sum",
34     fill_value=0
35 )
36
37 # Print matrix shape to confirm scale of recommendation system
38 print("Customer-Product Matrix Shape:", pivot_table.shape)
39
40 # -----
41 # Compute cosine similarity between products
42 # Similarity is calculated between product vectors (columns), so we transpose pivot_table
43 # -----
44 product_similarity = cosine_similarity(pivot_table.T)
45
46 # Convert similarity matrix into a DataFrame for easier lookup and recommendations
47 product_similarity_df = pd.DataFrame(
48     product_similarity,
49     index=pivot_table.columns,
50     columns=pivot_table.columns
51 )
52
53 # Print confirmation and similarity matrix size
54 print("Product similarity matrix created.")
55 print("Similarity Matrix Shape:", product_similarity_df.shape)
56
57 # -----
58 # Function: Recommend top N similar products based on cosine similarity
59 # This matches the project requirement of item-based collaborative filtering
60 # -----
61 def recommend_products(product_name, top_n=5):
62     # Handle case where product is not available in the dataset
63     if product_name not in product_similarity_df.index:
64         return f"Product '{product_name}' not found in dataset."
65
66     # Sort similar products by similarity score (highest similarity first)
67     similarity_scores = product_similarity_df[product_name].sort_values(ascending=False)
68
69     # Remove the same product itself from recommendations
70     similarity_scores = similarity_scores.drop(product_name)
71
72     # Return only the top N recommended product names
73     return list(similarity_scores.head(top_n).index)
74
75 # -----
76 # Test recommendation output using one sample product from the dataset
77 # -----
78 test_product = pivot_table.columns[8]
79 print("\nExample Product:", test_product)
80 print("Top 5 Recommendations:", recommend_products(test_product))
81
82 # Q1. Why do we create a CustomerID-Description-pivot-table for recommendations?
```

## 11. Hypothesis Testing (Statistical Validation)

### 11.1 Purpose of Hypothesis Testing in This Project

Hypothesis testing was conducted to statistically validate whether the observed patterns in segmentation and recommendations are meaningful and not random. This improves the credibility of the business insights and confirms that both the clustering and recommendation outputs can support real-world decision-making.

## 11.2 Hypothesis Test 1: Country Spending Difference (United Kingdom vs Germany)

**Objective:** To verify whether transaction spending behavior differs significantly between two major markets.

**Null Hypothesis (H0):** Transaction monetary values in the UK and Germany come from the same distribution (no difference).

**Alternate Hypothesis (H1):** Transaction monetary values differ significantly between the UK and Germany.

**Statistical Test Used:** Mann–Whitney U Test (two-sided)

**Result:**

- p-value = 2.35e-09 (< 0.05)

**Decision:** Reject H0

**Inference:** Transaction-level spending patterns vary significantly by country, meaning region-based pricing and promotional strategies can be optimized differently for each market.

## 11.3 Hypothesis Test 2: Monetary Spending Difference Across Customer Clusters

**Objective:** To validate that the customer segments created by clustering represent distinct spending behaviors.

**Null Hypothesis (H0):** All clusters have the same monetary distribution (no spending difference).

**Alternate Hypothesis (H1):** At least one cluster differs significantly in monetary behavior.

**Statistical Test Used:** Kruskal–Wallis H Test

**Result:**

- Test Statistic ≈ 1181.47

- p-value = 5.50e-248 (< 0.05)

**Decision:** Reject H0

**Inference:** Customer segments have statistically significant differences in spending behavior, confirming that segmentation is meaningful and suitable for targeted business strategies.

## 11.4 Hypothesis Test 3: Recommendation Quality vs Random Baseline

**Objective:** To verify that recommended products have significantly higher similarity than random product suggestions.

**Null Hypothesis (H0):** Top-5 recommendation similarity is not greater than random similarity.

**Alternate Hypothesis (H1):** Top-5 recommendations have significantly higher similarity than random products.

**Statistical Test Used:** Mann–Whitney U Test (one-sided, greater)

**Result:**

- **Top-5 similarity mean ≈ 0.8623**
- **Random similarity mean ≈ 0.0088**
- **p-value ≈ 0.0 (< 0.05)**  
**Decision:** Reject H<sub>0</sub>

**Inference:** The recommendation engine is statistically proven to outperform random suggestions, validating that item-based collaborative filtering produces meaningful personalization outputs.

## 11.5 Statistical Conclusions & Business Impact

All three hypothesis tests confirmed that segmentation patterns, geographic spending differences, and recommendation relevance are statistically significant. This strengthens the reliability of business insights and supports deployment-ready decision-making for personalization, retention campaigns, and market-level growth strategies.

```
LENOVO@DESKTOP-BRUVG0L MINGW64 ~/Desktop/ShopperSpectrum (main)
$ python hypothesis_test_1_country_revenue.py

Mann-Whitney U Test Results
Test Statistic: 3181776.5
P-value: 2.3469055229706595e-09

Conclusion: Reject H0
There is a statistically significant difference in transaction monetary values between UK and Germany.
((venv) )
```

```
LENOVO@DESKTOP-BRUVG0L MINGW64 ~/Desktop/ShopperSpectrum (main)
$ python hypothesis_test_1_country_revenue.py
((venv) )
LENOVO@DESKTOP-BRUVG0L MINGW64 ~/Desktop/ShopperSpectrum (main)
$ python hypothesis_test_2_cluster_monetary.py
● Clusters available: [np.int64(0), np.int64(1), np.int64(2), np.int64(3)]
Cluster 0 -> Customers: 3054, Avg Monetary: 1359.05
Cluster 1 -> Customers: 1067, Avg Monetary: 480.62
Cluster 2 -> Customers: 13, Avg Monetary: 127338.31
Cluster 3 -> Customers: 204, Avg Monetary: 12709.09

Kruskal-Wallis Test Results
Test Statistic: 1145.2656277243555
P-value: 5.501747703443845e-248

Conclusion: Reject H0
Customer segments have statistically significant differences in Monetary spending.
((venv) )
```

```
Kruskal-Wallis Test Results
Test Statistic: 1145.2656277243555
P-value: 5.501747703443845e-248

Conclusion: Reject H0
Customer segments have statistically significant differences in Monetary spending.
((venv) )
```

```

● $ python hypothesis_test_3_recommendation_quality.py
Running hypothesis test on recommendation similarity...

Top-5 similarity samples: 1000
Random similarity samples: 1000

Top-5 Similarity Mean: 0.581375791019578
Random Similarity Mean: 0.02806761074139078

Mann-Whitney U Test Results
Test Statistic: 988720.5
P-value: 0.0

Conclusion: Reject H0
The recommendation system produces significantly higher similarity than random suggestions.
((venv) )

```

```

● hypothesis_test_1_country_revenue.py > ...
1 import pandas as pd
2 from scipy.stats import mannwhitneyu
3
4 # Load cleaned dataset (fast + already processed)
5 df = pd.read_parquet("data/online_retail_cleaned.parquet")
6
7 # We test transaction-level monetary (InvoiceNo total amount)
8 txn_total = df.groupby(["InvoiceNo", "Country"])["TotalPrice"].sum().reset_index()
9
10 # Filter only two countries for comparison
11 uk = txn_total[txn_total["Country"] == "United Kingdom"]["TotalPrice"]
12 germany = txn_total[txn_total["Country"] == "Germany"]["TotalPrice"]
13
14 print("UK transactions:", len(uk))
15 print("Germany transactions:", len(germany))
16
17 # Mann-Whitney U Test (non-parametric)
18 # H0: UK and Germany transaction monetary values come from same distribution
19 # H1: They are different
20 stat, p_value = mannwhitneyu(uk, germany, alternative="two-sided")
21
22 print("\nMann-Whitney U Test Results")
23 print("Test Statistic:", stat)
24 print("P-value:", p_value)
25
26 # Conclusion (alpha = 0.05)
27 alpha = 0.05
28 if p_value < alpha:
29     print("\nConclusion: Reject H0")
30     print("There is a statistically significant difference in transaction monetary values between UK and Germany.")
31 else:
32     print("\nConclusion: Fail to Reject H0")
33     print("No statistically significant difference found between UK and Germany transaction monetary values.")
34

```

```

● hypothesis_test_2_cluster_monetary.py > ...
1 import pandas as pd
2 from scipy.stats import kruskal
3
4 # Load segmented customers (KMeans output)
5 segments_path = "data/customer_segments.csv"
6 rfm = pd.read_csv(segments_path)
7
8 # Ensure correct datatype
9 rfm["Cluster"] = rfm["Cluster"].astype(int)
10
11 print("Clusters available:", sorted(rfm["Cluster"].unique()))
12
13 # Prepare Monetary values for each cluster
14 cluster_groups = []
15 for c in sorted(rfm["Cluster"].unique()):
16     group = rfm[rfm["Cluster"] == c][["Monetary"]]
17     cluster_groups.append(group)
18     print(f"\nCluster {c} -> Customers: {len(group)}, Avg Monetary: {(group.mean()):.2f}")
19
20 # Kruskal-Wallis test (non-parametric)
21 # H0: All clusters have the same spending distribution
22 # H1: At least one cluster differs
23 stat, p_value = kruskal(*cluster_groups)
24
25 print("\nKruskal-Wallis Test Results")
26 print("Test Statistic:", stat)
27 print("P-value:", p_value)
28
29 # Conclusion (alpha = 0.05)
30 alpha = 0.05
31 if p_value < alpha:
32     print("\nConclusion: Reject H0")
33     print("Customer segments have statistically significant differences in Monetary spending.")
34 else:
35     print("\nConclusion: Fail to Reject H0")
36     print("No statistically significant difference found between cluster Monetary spending.")
37

```

```

● hypothesis_test_3_recommendation_quality.py > ...
28 # Pick random sample of products to evaluate
29 sample_products = np.random.choice(product_list, size=min(num_trials, len(product_list)), replace=False)
30
31 for product in sample_products:
32     # Get similarity scores for the chosen product
33     sim_scores = product_similarity_df[product].sort_values(ascending=False)
34
35     # Remove self-similarity
36     sim_scores = sim_scores.drop(product)
37
38     # Top 5 recommendation similarities
39     top5 = sim_scores.head(top_n).values
40     top5_scores.extend(top5)
41
42     # Random 5 product similarities (baseline)
43     random_items = np.random.choice(sim_scores.index, size=top_n, replace=False)
44     random_vals = sim_scores.loc[random_items].values
45     random_scores.extend(random_vals)
46
47 top5_scores = np.array(top5_scores)
48 random_scores = np.array(random_scores)
49
50 print("Top-5 similarity samples:", len(top5_scores))
51 print("Random similarity samples:", len(random_scores))
52
53 print("\nTop-5 Similarity Mean:", top5_scores.mean())
54 print("Random Similarity Mean:", random_scores.mean())
55
56 # Mann-Whitney U test (one-sided)
57 # H0: top5_scores <= random_scores
58 # H1: top5_scores > random_scores
59 stat, p_value = mannwhitneyu(top5_scores, random_scores, alternative="greater")
60
61 print("\nMann-Whitney U Test Results")
62 print("Test Statistic:", stat)
63 print("P-value:", p_value)
64
65 alpha = 0.05
66 if p_value < alpha:
67     print("\nConclusion: Reject H0")
68     print("The recommendation system produces significantly higher similarity than random suggestions.")
69 else:
70     print("\nConclusion: Fail to Reject H0")
71     print("The recommendation system does not significantly outperform random product suggestions.")
72

```

## 12. Model Deployment & Streamlit Application

### 12.1 Streamlit UI Overview

A multi-page Streamlit application was developed to deploy the project results in an interactive and user-friendly format. The application provides real-time outputs for customer segmentation and product recommendations, allowing business users to explore insights without manual data analysis.

### 12.2 Customer Segmentation Module (RFM → Segment Prediction)

This module accepts three customer behavior inputs:

- **Recency (days)**
- **Frequency (number of purchases)**
- **Monetary (total spend)**

The saved preprocessing and model artifacts were reused for prediction:

- models/scaler.pkl
- models/kmeans\_model\_k4.pkl
- models/segment\_map.json

**Output:** Predicted cluster and segment label (High Value / Regular / Occasional / At Risk), along with an actionable business recommendation.

### 12.3 Product Recommendation Module (Product Input → Top 5 Similar Products)

This module allows users to select a product name and instantly receive **5 similar product recommendations** based on cosine similarity. It uses saved recommendation artifacts:

- models/product\_similarity.pkl
- models/product\_list.pkl

**Output:** Top 5 recommended product names displayed in a clean card-style UI for better usability and presentation.

### 12.4 Business Insights Dashboard (KPIs + Segment Distribution + Use Cases)

The dashboard module presents core business indicators computed from the cleaned dataset and labeled segments:

- **Total Customers** (based on **CustomerID**)
- **Total Transactions** (based on unique **InvoiceNo**)

- **Unique Products** (based on unique Description)
- **Total Revenue** (sum of TotalPrice)

It also displays customer segment distribution and explains real-time business use cases such as retention, targeted marketing, dynamic pricing, and inventory planning.

## 12.5 Model Saving & Loading Strategy (Joblib + JSON)

To ensure real-time prediction performance and consistent results across runs, models and artifacts were stored using:

- **Joblib (.pkl)** for trained models and recommendation data
- **JSON** for segment mapping labels

```
◆ app.py > ...
  # It increases adoption and makes the system feel like a deployable analytics product.
13
14 import streamlit as st
15
16 # Configure the Streamlit page settings such as title, icon, and layout width
17 st.set_page_config(
18     page_title="Shopper Spectrum",
19     page_icon="🛍️",
20     layout="wide"
21 )
22
23 # Adding custom CSS styling to make the homepage look more professional
24 # This controls font sizes, colors, and card design for a clean UI/UX
25 st.markdown(
26     """
27     <style>
28         .main-title {
29             font-size: 44px;
30             font-weight: 800;
31             margin-bottom: 5px;
32         }
33         .sub-title {
34             font-size: 18px;
35             color: #667788;
36             margin-top: 0px;
37         }
38         .card {
39             background-color: #ffffff;
40             padding: 18px;
41             border-radius: 14px;
42             border: 1px solid #e6e7e8;
43             box-shadow: 0px 2px 8px rgba(0,0,0,0.05);
44         }
45     </style>
46     """
47     unsafe_allow_html=True
48 )
49
50 # Main title and subtitle for the homepage
51 st.markdown(<div class="main-title">Shopper Spectrum</div>, unsafe_allow_html=True)
52 st.markdown(<div class="sub-title">Customer Segmentation and Product Recommendation System</div>, unsafe_allow_html=True)
53
54 # Adding spacing for better layout
55 st.write("")
56 st.write("")
57
58 # Creating a 3-column layout to highlight the main modules of the project
59 col1, col2, col3 = st.columns(3)
60
61 # Adding spacing for better layout
62 st.write("")
```

```
63     # Card 1: Customer Segmentation module overview
64     st.markdown(
65         """
66         <div class="card">
67             <h3>Customer Segmentation</h3>
68             <p>Segment customers using RFM analysis and KMeans clustering for targeted marketing strategies.</p>
69         </div>
70     """
71     unsafe_allow_html=True
72 )
73
74 with col2:
75     # Card 2: Product Recommendation module overview
76     st.markdown(
77         """
78         <div class="card">
79             <h3>Product Recommendation</h3>
80             <p>Recommend 5 similar products using item-based collaborative filtering and cosine similarity.</p>
81         </div>
82     """
83     unsafe_allow_html=True
84 )
85
86 with col3:
87     # Card 3: Business Insights module overview
88     st.markdown(
89         """
90         <div class="card">
91             <h3>Business Insights</h3>
92             <p>Enable retention strategies, dynamic pricing, and inventory optimization using customer behavior insights.</p>
93         </div>
94     """
95     unsafe_allow_html=True
96 )
97
98 # Extra spacing and navigation guidance for the user
99 st.write("")
100 st.info("Use the sidebar to navigate: Customer Segmentation and Product Recommendation.")
```

```
pages > ◆ 1.Customer_Segmentation.py > ...
20 import json
21
22 # Page configuration for better layout and page title
23 st.set_page_config(page_title="Customer Segmentation", layout="wide")
24
25 # Page heading and short description
26 st.title("Customer Segmentation")
27 st.write("Predict customer segment using RFM values (Recency, Frequency, Monetary).")
28
29 #
30 # Loading the saved ML models (trained earlier)
31 # This avoids retraining the model every time the app runs
32 #
33 scaler = joblib.load("models/scaler.pkl")
34 kmeans = joblib.load("models/kmeans_model_k4.pkl")
35
36 #
37 # Loading segment mapping (Cluster ID -> Segment Name)
38 # This helps show meaningful labels instead of just numbers
39 #
40 with open("models/segment_map.json", "r") as f:
41     segment_map = json.load(f)
42
43 # Convert JSON keys (strings) into integers for correct mapping
44 segment_map = {int(k): v for k, v in segment_map.items()}
```

```
pages > ◆ 1.Customer_Segmentation.py > ...
61     )
62
63 frequency = st.number_input(
64     "frequency (number of purchases)",
65     min_value=0,
66     value=5,
67     step=1
68 )
69
70 monetary = st.number_input(
71     "monetary (total spend)",
72     min_value=0,
73     value=1000.0,
74     step=10.0
75 )
76
77 # Button triggers the prediction
78 Predict_BTN = st.button("Predict Segment")
79
80 #
81 # Right side: Prediction Output Section
82 #
83 with col2:
84     pages > ◆ 3.Business_Insights.py > ...
85
86     # Q4. Why do we include business use cases like dynamic pricing and inventory optimization in the insights?
87     # Answer: It proves that the analytics pipeline is not just technical, but directly supports revenue, growth, and operational efficiency.
88     # These use cases convert customer behavior patterns into practical strategies that stakeholders can implement.
89
90     import streamlit as st
91     import pandas as pd
92
93     # Page configuration for better layout and page title
94     st.set_page_config(page_title="Business Insights", layout="wide")
95
96     # Page heading and short description
97     st.title("Business Insights Dashboard")
98     st.write("Key insights and real-time use cases derived from customer behavior and purchase patterns.")
99
100     # Divider for better structure
101     st.markdown("----")
102
103     # -----
104     # Load cleaned transaction data and labeled customer segments
105     # Using cleaned parquet improves speed and ensures valid transactions
106     # -----
107     df = pd.read_parquet("data/online_retail_cleaned.parquet")
```

```
pages > ◆ 2.Product_Recommendation.py > ...
10 st.title("Product Recommendation")
11 st.write("Select a product name and get 5 similar recommendations based on cosine similarity.")
12
13 #
14 # Load precomputed similarity matrix and product list
15 # This avoids expensive computation inside Streamlit
16 #
17 product_similarity_df = joblib.load("models/product_similarity.pkl")
18 product_list = joblib.load("models/product_list.pkl")
19
20 #
21 # Custom CSS for card-style recommendation output
22 # Improves UI/UX and makes recommendations look professional
23 #
24 st.markdown(
25     """
26     <style>
27         .rec-card {
28             background-color: #ffffff;
29             padding: 14px;
30             border-radius: 12px;
31             border: 1px solid #e6e7e8;
32             margin-bottom: 10px;
33             box-shadow: 0px 2px 8px rgba(0,0,0,0.05);
34             font-size: 16px;
35             font-weight: 600;
36         }
37     </style>
38     """
39     unsafe_allow_html=True
40 )
41
42 # Convert JSON keys (strings) into integers for correct mapping
43 segment_map = {int(k): v for k, v in segment_map.items()}
```

<<

Deploy :

app

Customer Segmentation

Product Recommendation

Business Insights

# Shopper Spectrum

Customer Segmentation and Product Recommendation System

## Customer Segmentation

Segment customers using RFM analysis and KMeans clustering for targeted marketing strategies.

## Product Recommendation

Recommend 5 similar products using item-based collaborative filtering and cosine similarity.

## Business Insights

Enable retention strategies, dynamic pricing, and inventory optimization using customer behavior insights.

Use the sidebar to navigate: Customer Segmentation and Product Recommendation.

# Customer Segmentation

Predict customer segment using RFM values (Recency, Frequency, Monetary).

Recency (days since last purchase)

12

- +

## Prediction Output

Frequency (number of purchases)

15

- +

Predicted Cluster: 3

Monetary (total spend)

10500.00

- +

Customer Segment: Regular

Predict Segment

Recommended Action: Cross-sell bundles and membership benefits to increase repeat purchases.

# Customer Segmentation

Predict customer segment using RFM values (Recency, Frequency, Monetary).

Recency (days since last purchase)

3

- +

Frequency (number of purchases)

11

- +

Monetary (total spend)

10500.00

- +

Predict Segment

## Prediction Output

Predicted Cluster: 0

Customer Segment: Occasional

Recommended Action: Offer discounts and personalized recommendations to increase engagement.

# Product Recommendation

Select a product name and get 5 similar recommendations based on cosine similarity.

Select Product Name

12 HANGING EGGS HAND PAINTED

## Top 5 Recommended Products

1. SET OF 6 EASTER RAINBOW CHICKS

2. BAG OF SILVER STONES

3. UBO-LIGHT TRIOBASE BLUE

4. ORIENTAL RED CUSHION COVER

5. GREEN PEONY CUSHION COVER

# Product Recommendation

Select a product name and get 5 similar recommendations based on cosine similarity.

Select Product Name

3 PIECE SPACEBOY COOKIE CUTTER SET

Get Recommendations

## Top 5 Recommended Products

1. MEMO BOARD COTTAGE DESIGN

2. SET 12 COLOURING PENCILS DOILEY

3. SET 36 COLOURING PENCILS DOILEY

4. I LOVE LONDON MINI RUCKSACK

5. DOLLCRAFT GIRL AMELIE KIT

# Business Insights Dashboard

Key insights and real-time use cases derived from customer behavior and purchase patterns.

Total Customers	Total Transactions	Unique Products	Total Revenue
4,338	18,532	3,866	8,911,407.90

## Customer Segment Distribution

	Segment	Customers	
0	Occasional	3054	3054
1	At Risk	1067	1067
2	Regular	204	204
3	High Value	13	13

## Real-time Business Use Cases

### Customer Segmentation for Targeted Marketing Campaigns

- High Value: premium loyalty rewards, exclusive offers, personalized bundles
- Regular: cross-sell campaigns, membership benefits, repeat-purchase reminders
- Occasional: discounts, trending products, engagement campaigns
- At Risk: win-back offers, reminders, limited-time coupons

### Personalized Product Recommendations on E-Commerce Platforms

- Recommend 5 similar products using cosine similarity
- Improves customer experience and increases average order value
- Helps cross-selling and product discovery

### Identifying At-Risk Customers for Retention Programs

- At Risk customers have high recency and low frequency
- Retention actions: reminder emails, personalized discounts, win-back campaigns
- Prevents churn and improves repeat purchases

### Dynamic Pricing Strategies Based on Purchase Behavior

- High demand products can maintain price with minimal discounts
- Low demand products can be promoted with discounts to increase sales
- Segment-based offers improve revenue without over-discounting

# 13. Business Insights Dashboard (Final Outputs)

## 13.1 Key Business KPIs

The dashboard summarizes key operational metrics that reflect overall business health and performance. These KPIs include total unique customers, total invoice transactions, unique product count, and total revenue derived from all valid purchases using the engineered **TotalPrice** feature.

## 13.2 Segment Distribution Summary (With Numeric Results)

Customer segmentation revealed a strong imbalance in customer value contribution, which is common in retail. The final segment distribution was:

- **Occasional Customers:** 3054
- **At Risk Customers:** 1067
- **Regular Customers:** 204
- **High Value Customers:** 13

**Insight:** Although High Value customers are the smallest group, they represent the highest revenue opportunity, while Occasional customers represent the largest conversion opportunity by volume.

## 13.3 Revenue Contribution Patterns (Pareto Insight)

Revenue distribution analysis confirmed that customer spend is highly skewed, where a small fraction of customers contribute disproportionately to total revenue. This validates the need to retain and prioritize high-value and regular customers through loyalty strategies rather than applying uniform discounts across the customer bases.

## 13.4 Actionable Recommendations per Segment

Each segment enables direct business decision-making:

- **High Value (13 customers):** premium rewards, exclusive access offers, priority support
- **Regular (204 customers):** bundle offers, cross-sell recommendations, membership upgrades
- **Occasional (3054 customers):** discounts, engagement campaigns, trending product discovery
- **At Risk (1067 customers):** win-back offers, reminders, time-limited incentives

**Inference:** Segment-based targeting improves campaign efficiency by aligning marketing actions with customer profitability and churn risk.

## 13.5 Decision-Making Impact (Marketing, Retention, Inventory)

The dashboard connects analytics to business strategy by enabling:

- smarter allocation of marketing budget across segments
- early identification of churn-prone customers for retention intervention
- inventory prioritization based on high-frequency customer demand signals
- product recommendation-driven cross-selling to improve average order value

# 14. Project Constraints & Practical Considerations

## 14.1 Handling Large Dataset Efficiently

The raw dataset contained **541,909 rows**, making repeated CSV loading and heavy pivot computations expensive for a local system. To optimize performance, the dataset was cleaned once and converted into a **Parquet file**, reducing load time and enabling faster downstream processing for EDA, RFM generation, and model training.

## 14.2 Local Machine Limitations (RAM/Storage Constraints)

The project was executed on a system with **8 GB RAM** and limited storage availability (~8 GB free). This required efficient memory handling and avoiding unnecessary recomputation, especially during recommendation matrix creation and visualization generation.

## 14.3 Recommendation Matrix Complexity & Scalability

The recommendation model required building a customer-product interaction matrix and computing cosine similarity across products. The pivot table size was:

- **4338 customers × 3866 products**

This means product-to-product similarity can scale to millions of similarity pairs, making it computationally heavy. To address this, similarity results were saved using Joblib so that recommendations load instantly inside Streamlit without recomputation.

## 14.4 Model Output Imbalance and Business Interpretability

Model comparison revealed that DBSCAN produced highly imbalanced clustering results:

- **4262 customers in one cluster, 9 in another, and 67 noise points**

Even though DBSCAN gave a high silhouette score (**0.6662**), it was not suitable for segmentation-driven marketing because it could not generate usable customer groups. KMeans was selected because it produced stable and interpretable segments aligned with business requirements.

# 15. Final Conclusion

## 15.1 Summary of Achievements

The project successfully transformed raw retail transaction data (**541,909 rows**) into a complete analytics pipeline including data cleaning, EDA, RFM profiling, customer segmentation, product recommendations, statistical validation, and real-time Streamlit deployment. The final outputs are both technically solid and directly aligned with real-world business decision-making.

## 15.2 Business Impact Delivered (With Numeric Results)

Customer segmentation created four business-ready segments with clear behavior patterns:

- **Occasional:** 3054 customers
- **At Risk:** 1067 customers
- **Regular:** 204 customers
- **High Value:** 13 customers

This distribution confirms that the business has a large conversion opportunity in the Occasional group and a major revenue-protection requirement for the small but extremely valuable High Value group.

## 15.3 What Worked Best & Why (Model + Recommendation Validation)

Among three clustering models (KMeans, Agglomerative, DBSCAN), **KMeans (k=4)** was selected because it provided the most balanced combination of clustering performance and business interpretability (**Silhouette = 0.6162, DBI = 0.7534**). The recommendation system also proved statistically meaningful with **Top-5 similarity mean ≈ 0.8623**, compared to a random baseline ( $\approx 0.0088$ ), confirming strong personalization quality.

## 15.4 Final Recommendations for Stakeholders

The business should prioritize retention of High Value and Regular customers through loyalty incentives and premium offers, while using recommendations and discounts to increase frequency among Occasional customers. At Risk customers should be targeted with win-back campaigns, reminders, and time-based offers to reduce churn and maximize customer lifetime value.